

CS229 Notes: Regularization & Model Selection

Section 1: Cross-Validation $\rightarrow$

Core Problem :-

Several models? $\rightarrow$ which one to pick?

~~Wrong~~ instinct: pick the one with lowest training error.
Why wrong? $\rightarrow$ A 10-degree polynomial will memorize to point perfectly.
 $\rightarrow$ Training error = 0
but, fails on new data.
 $\rightarrow$ high variance/overfitting.

Actually we should choose a model based on ~~error~~ generalization error which is during prediction using testing.

Soln $\rightarrow$ use $\rightarrow$

(i) Hold-Out Cross Validation :-

(split data) $\rightarrow$

70% $\rightarrow$ training (S-train)

30% $\rightarrow$ test (S-test) (model never sees ~~data~~ during training)

Run $\rightarrow$ 1000 house price examples

• M1: linear fit (degree 1)

• M2: quadratic (2nd degree)

• M5: 5 degree

Train all three

• M1: 15% error

• M2: 9% error

• M5: 22% error (overfitting kicks in)

on 700 examples $\rightarrow$ Test error on 300 held-out:

$\rightarrow$ picking M2

that we wasted 300 examples.

(ii) Soln $\rightarrow$

• K-Fold Cross Validation :-

Soln $\rightarrow$ reuse the data smarter.

$\rightarrow$ split into $k=10$ equal chunks. For each chunk j :

• Train on the other 9 chunks

• Test on chunk $j \rightarrow$ could be 7% / 8% / 5% / 1%

• get error $\rightarrow$ on j


Run $\rightarrow$ (concrete example $\rightarrow k=3, 9 \Rightarrow$

Fold 1: train on [4, 5, 6, 7, 8, 9], test on $\rightarrow$ [1, 2, 3] $\rightarrow$ error = 12%.

Fold 2: train $\rightarrow$ [1, 2, 3, 7, 8, 9], test on $\rightarrow$ [4, 5, 6] $\rightarrow$ error = 10%.

Fold 3: train $\rightarrow$ [1, 2, 3, 4, 5, 6], test $\rightarrow$ [7, 8, 9] $\rightarrow$ error = 19%.

Estimated generalization error = $(12 + 10 + 19) / 3 = 12\%$



How are we going to identify which model is sensitive to small data changes and which are to large data changes without doing experiment testing (costly)?

Model known to be sensitive for small data changes →

- ① Neural Networks — different random initialization & slightly different data = very different weights.
- ② Decision trees → diff point near root → diff tree structure
- ③ Boosting methods → early weak learning affect everything downstream

Models that are relatively stable →

- ① Linear/Logistic Regression — convex loss, unique sol, small data changes = small parameter changes.
- ② Naive Bayes — counting probabilities, very stable.
- ③ SVM with large margin — support vectors dominate, most points don't even matter.

Rule: — If model has a non-convex loss or involves sequentially/greedy decisions → be cautious about retraining. → Sensitive

If it has a convex loss with a unique sol → retrain freely. → non sensitive

Sec-3: Bayesian statistics & Regularization

in M.L.E → "what θ makes the data most likely?"

$$\theta_{M.L.E} = \arg \max_{\theta} \prod_{i=1}^n P(x^{(i)} | \mu^{(i)}; \theta)$$

Assumption → θ is fixed unknown constant, it exist, we just don't know it.

They Bayesian Shift $\rightarrow$ says $\rightarrow \theta$ is not a fixed constant. It is random variable with its own distribution. Before seeing any data, we have prior belief about θ $\rightarrow$ written $p(\theta)$.

Analogy $\Rightarrow$ Guessing std exam score $= \theta$,
 without prior, scores $\rightarrow$ 40-90, clustered around 70 (clumpy)
 with prior, homework grades (data) $\rightarrow$ updating belief

The posterior after seeing training data S $\rightarrow$ updated belief $\rightarrow$ posterior

$$p(\phi|S) = \underbrace{p(S|\phi)}_{\text{likelihood}} \cdot \underbrace{p(\phi)}_{\text{prior}} \cdot \underbrace{\frac{1}{p(S)}}_{\text{normalizing constant (can be ignored)}}$$

[posterior = likelihood x prior]

Fully Bayesian prediction $\Rightarrow$

To pred on x (new data) $\xrightarrow{\text{ans}}$ all possible θ , weighted by how probable each θ is:

$$p(y|x, S) = \int_{\theta} p(y|x, \theta) p(\theta|S) d\theta$$

all possible θ how probable each θ is

Intuition $\Rightarrow$ don't commit to 1 θ , consider all plausible θ values, (make sense)

Problem $\Rightarrow$ Integral is impossible to compute in closed form (θ is high dimensional - imagine integrating over 10,000 dim).

MAP: - The practical Approximation.
 Instead of integrating over all θ , just find the single most probable θ under posterior $\Rightarrow$

$$\theta_{\text{MAP}} = \arg \max_{\theta} \prod_{i=1}^n p(y^{(i)}|x^{(i)}; \theta) \cdot p(\theta)$$

closed form v/s other sol's
 $\downarrow$
 exact compact self-tuning formula ($\propto$)
 $\downarrow$
 numerical approx (e.g. 1.415)

Compare with ML: $\Rightarrow$

$$\theta_{\text{ML}} = \arg \max_{\theta} \prod_{i=1}^n p(y^{(i)}|x^{(i)}; \theta)$$

closed form v/s open form

$\downarrow$
 finite ops ($t, x, 1, \log, \exp$)
 $\downarrow$
 infinite no of operations
 $\downarrow$
 no iterations or loops (hard to code)
 $\downarrow$
 infinite slices of parameter space
 $\downarrow$
 functions used (polynomial, exponential, trigonometric, etc.)

$\rightarrow$ Diff is MAP has extra $p(\theta)$
 probable values of θ

Now Connection to Regularization:-

what if your prior is $\Rightarrow \theta \sim N(0, \sigma^2 I)$

"I believe θ is probably small-centered at zero"

Taking log of the MAP objective (log turns products $\rightarrow$ sum)

$$\theta_{\text{MAP}} = \arg \max_{\theta} \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}, \theta) + \log p(\theta)$$

• Plug in the Gaussian prior - log of Gaussian is a quadratic:

$$\log p(\theta) \propto -\frac{1}{2\sigma^2} \|\theta\|^2$$

• so MAP becomes:

$$\theta_{\text{MAP}} = \arg \max_{\theta} \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}, \theta) - \frac{1}{2\sigma^2} \|\theta\|^2$$

• Equivalently (flipping sign, max $\rightarrow$ min):

$$\theta_{\text{MAP}} = \arg \min_{\theta} \underbrace{\sum \text{loss}}_{\text{fit the data}} + \underbrace{\lambda \|\theta\|^2}_{\text{keep } \theta \text{ small}}$$

we found L_2 regularization (Ridge Regression)

• The Great Intuition:-

MLE: "Find θ that explains the data best.
I don't care how large θ gets."

MAP with Gaussian prior: "Find θ that explains the data well, BUT I have prior belief θ should be small. Don't go crazy with large weights."

Large weights = model is very sensitive to i/p phases = Overfitting that. The prior penalizes

this is what regularization is. $\lambda ||\theta||^2$

• Why MAP and Regularization reduces overfitting?

⇒ BEZ of prior term $\rightarrow p(\theta)$ in MLE $\rightarrow$ similar in regularization